

Development of a Serverless Web Application with AWS

Academic Report

Prepared by:

Kenza Elhamchi
Zakaria Benlamkadam

Supervised by:

Hayat Routaib

December 2024

Contents

| | | |
|----------|---|-----------|
| 1 | Introduction | 2 |
| 1.1 | Project Context and Objective | 2 |
| 1.2 | Importance of Serverless Architectures | 2 |
| 1.3 | Overview of Tools Used | 2 |
| 2 | Project Description | 3 |
| 2.1 | Project Objective | 3 |
| 2.1.1 | Problems Addressed: | 3 |
| 2.2 | Use Cases | 3 |
| 2.3 | System Architecture | 3 |
| 2.3.1 | Overview | 3 |
| 2.3.2 | Role of Each Component | 4 |
| 3 | Technical Implementation | 5 |
| 3.1 | 1. Frontend Development | 5 |
| 3.1.1 | 1.1 Website Structure | 5 |
| 3.1.2 | 1.2 Description of HTML, CSS, and JavaScript Files Used | 5 |
| 3.2 | 2. Backend Development | 6 |
| 3.2.1 | 2.1 Lambda Function Definitions | 6 |
| 3.2.2 | 2.2 Description of Interactions with DynamoDB | 7 |
| 3.3 | 3. AWS Configuration | 7 |
| 3.3.1 | 3.1 Steps to Configure S3 (Static Hosting) | 7 |
| 3.3.2 | 3.2 Steps to Configure CloudFront | 8 |
| 3.3.3 | 3.3 Creating APIs with API Gateway | 10 |
| 3.3.4 | 3.4 Deploying Lambda Functions | 11 |
| 3.3.5 | 3.5 Configuring DynamoDB (Table and Data Schema) | 13 |
| 4 | Data Flow | 14 |
| 4.0.1 | 4.1 Step-by-Step Explanation of a User Request | 14 |
| 5 | Project Benefits and Limitations | 14 |
| 5.1 | 4.1 Benefits | 14 |
| 5.2 | 4.2 Limitations | 15 |
| 6 | Results | 15 |
| 7 | Conclusion and Future Work | 16 |

1 Introduction

1.1 Project Context and Objective

Modern businesses aim to develop web solutions that are cost-effective, high-performing, and scalable. Serverless architectures offer an attractive alternative to traditional systems, where the cloud provider manages infrastructure, scaling, and resource allocation automatically. This eliminates the need for server management and reduces infrastructure costs, allowing businesses to focus on development rather than operations.

The objective of this project is to build a web application for recording and retrieving student data using AWS serverless technologies. The solution is designed to be cost-efficient, scalable, and easy to maintain.

1.2 Importance of Serverless Architectures

Serverless architecture simplifies development by removing the complexity of infrastructure management. Key benefits include:

- **Elimination of Server Management:** No need for physical or virtual server management; infrastructure is handled by the cloud provider.
- **Pay-as-you-go Billing:** Costs are based on actual resource usage, reducing unnecessary expenses.
- **Automatic Scalability:** Resources automatically scale based on workload demands.
- **Rapid Development:** Services like AWS Lambda reduce development time with simplified configurations.

These advantages make serverless a perfect fit for applications requiring flexibility and cost-effectiveness.

1.3 Overview of Tools Used

The project uses the following AWS services:

- **Amazon S3:** Hosts the frontend (static files like HTML, CSS, and JavaScript).
- **Amazon CloudFront:** Distributes the static content hosted on S3 globally, reducing latency by caching the content at edge locations closer to users, ensuring faster load times and improved performance.
- **AWS Lambda:** Executes backend code in response to API requests without managing servers.
- **Amazon API Gateway:** Manages HTTP requests to Lambda functions.
- **Amazon DynamoDB:** A NoSQL database for storing student data.

By integrating these tools, we developed a serverless web application that is scalable, cost-efficient, and easy to maintain.

2 Project Description

2.1 Project Objective

This project demonstrates how serverless architectures can address the scalability, cost, and management challenges in modern web applications. The main focus is on simplifying the management of student data while minimizing infrastructure costs. By using a serverless setup, this project eliminates backend infrastructure concerns and ensures seamless user experiences even during traffic surges.

2.1.1 Problems Addressed:

- **Server Management Complexity:** Traditional systems require constant monitoring and management of infrastructure, which is resolved with serverless services.
- **Scaling Challenges:** Traditional systems may not scale efficiently during traffic spikes. The serverless architecture automatically adapts to demand.
- **Data Management Efficiency:** Serverless solutions, like DynamoDB, provide low-latency access to data and scale automatically.

2.2 Use Cases

The application includes the following core features:

- **Storing Student Data:** Users submit student details (e.g., name, class), which are processed and saved in DynamoDB through Lambda functions.
- **Retrieving Student Data:** Users can search for stored student information by providing an identifier or search parameter.
- **Optional Data Modification:** Features to update or delete student data may be added in future iterations.
- **Frontend Hosting:** The user interface is hosted on AWS S3 as a static website.
- **Request Management:** API Gateway handles secure requests, with optional authentication mechanisms.

2.3 System Architecture

2.3.1 Overview

Global Architecture Diagram

Figure 20 illustrates the architecture of the system, showcasing the flow of data between components:

- **Frontend:** Hosted on AWS S3 and delivered through Amazon CloudFront. CloudFront caches static content (HTML, CSS, JavaScript, etc.) at edge locations to ensure faster load times and improved performance globally.
- **API Gateway:** Manages incoming HTTP requests.

- **Lambda:** Handles business logic and interacts with DynamoDB.
- **DynamoDB:** Stores and retrieves data.




Figure 1: Global Architecture Diagram

Explanation of the Main Components:

The serverless architecture leverages AWS services to minimize infrastructure management and ensure scalability:

- **Frontend:** Hosted as static files in an S3 bucket and delivered through Amazon CloudFront.
- **Backend:** The backend is serverless, with requests routed via API Gateway to Lambda functions, which then interact with DynamoDB.
- **Database:** DynamoDB stores application data efficiently, ensuring fast reads and writes.

2.3.2 Role of Each Component

S3 and CloudFront: Hosting the Application

- S3 stores the frontend (HTML, CSS, JavaScript).
- CloudFront is used to distribute the static content globally, caching it at edge locations to reduce latency and improve load times for end users.
- CloudFront serves the static files from S3 to the browser, ensuring high availability and minimal latency.

Benefits:

- Scales automatically and charges only for storage and data transfer. CloudFront further optimizes the delivery of static content by reducing latency and speeding up access for global users.

API Gateway: Managing HTTP Requests

- Acts as the entry point for HTTP requests from the frontend.
- Routes requests to Lambda functions based on the HTTP method (GET, POST, PUT, DELETE).

Benefits:

- Secure communication, request validation, and monitoring features are available.

Lambda: Executing Backend Code

- Lambda functions execute the application logic, such as validating inputs and interacting with DynamoDB.
- They are triggered by requests via API Gateway.

Benefits:

- Pay-per-use, automatic scaling, and high availability.

DynamoDB: Storing Data

- A NoSQL database for storing student records.

Features:

- High throughput, low latency, and automatic scaling for large datasets. Security is ensured with encryption at rest.

3 Technical Implementation

3.1 1. Frontend Development

The frontend of the application is developed using HTML, CSS, and JavaScript files. The goal is to create a simple and interactive user interface that allows interaction with the backend by submitting and retrieving data.

3.1.1 1.1 Website Structure

The user interface is organized to be clear and user-friendly. Here is the website structure:

- **index.html**: This is the main file that contains the structure of the homepage, including input fields for student information and buttons to interact with the API.
- **scripts.js**: This JavaScript file is essential for managing interactions between the frontend and backend. It allows sending HTTP requests to APIs and receiving responses.

3.1.2 1.2 Description of HTML, CSS, and JavaScript Files Used

- **HTML (index.html)**: The HTML file contains forms for submitting student information and a button to fetch data. It is designed to be responsive and suitable for different screen sizes. the css included inside the index.html
- **JavaScript (scripts.js)**: This file manages the logic for API calls. For example, it uses `fetch()` to send POST and GET requests to the API Gateway.

Example Code for API Call in scripts.js:

```
// AJAX POST request to save student data
document.getElementById("savestudent").onclick = function(){
    var inputData = {
        "name": $('#name').val(),
        "class": $('#class').val(),
        "age": $('#age').val()
    };

    $.ajax({
        url: API_ENDPOINT,
        type: 'POST',
        data: JSON.stringify(inputData),
        contentType: 'application/json; charset=utf-8',
        success: function (response) {

            let responseData = typeof response.body === 'string' ? JSON.parse(response.body) : response.body;

            document.getElementById("studentSaved").innerHTML = responseData.message || "Student Data Saved!";

        }
    });
};
```

Figure 2: Example Code for API Call in scripts.js

In this code, the `addStudent()` function sends a POST request to the API Gateway to save a student's data in DynamoDB.

Image to Insert: Screenshot of the JavaScript code (scripts.js).

3.2 2. Backend Development

The backend of the application is fully managed by AWS Lambda functions, enabling automatic resource management without server infrastructure concerns. Each Lambda function is triggered by HTTP calls through the API Gateway.

3.2.1 2.1 Lambda Function Definitions

The two main Lambda functions used in this project are:

- **GET Function:** This function retrieves a student's information based on their ID.

Example Code for GET Function:

```
lambda_function.py X
lambda_function.py
1  import json
2  import boto3
3
4  def lambda_handler(event, context):
5      # Initialize a DynamoDB resource object for the specified region
6      dynamodb = boto3.resource('dynamodb', region_name='eu-north-1')
7
8      # Select the DynamoDB table named 'studentData'
9      table = dynamodb.Table('StudentData')
10
11     # Scan the table to retrieve all items
12     response = table.scan()
13     data = response['Items']
14
15     # If there are more items to scan, continue scanning until all items are retrieved
16     while 'LastEvaluatedKey' in response:
17         response = table.scan(ExclusiveStartKey=response['LastEvaluatedKey'])
18         data.extend(response['Items'])
19
20     # Return the retrieved data
21     return data
22
```

Figure 3: Example Code for GET Function

- **POST Function:** This function saves new data in DynamoDB.

Example Code for POST Function:



```

23 def lambda_handler(event, context):
24     try:
25         # Obtenez un nouvel ID unique
26         student_id = get_next_id()
27
28         # Extraire les autres données
29         name = event.get('name')
30         student_class = event.get('class')
31         age = event.get('age')
32
33         # Insérer les données de l'étudiant
34         response = student_table.put_item(
35             Item={
36                 'studentId': str(student_id),
37                 'name': name,
38                 'class': student_class,
39                 'age': age
40             }
41         )
42
43         # Réponse JSON structurée
44         return {
45             'statusCode': 200,
46             'body': json.dumps({'message': 'Student data saved successfully!'})
47         }
48
49     except Exception as e:
50         print(f"Erreur : {e}")
51         return {
52             'statusCode': 500,
53             'body': json.dumps({'message': 'Erreur interne du serveur'})
54         }

```

Figure 4: Example Code for POST Function

3.2.2 2.2 Description of Interactions with DynamoDB

DynamoDB is used to store student information. Each student is identified by a unique `student_id`. When the POST function is called, the data is sent to DynamoDB where a new item is added. When the GET function is called, DynamoDB searches for the student corresponding to the provided ID and returns the relevant information.

3.3 3. AWS Configuration

This section details the steps to configure the necessary AWS services for this project: S3, API Gateway, Lambda, and DynamoDB.

3.3.1 3.1 Steps to Configure S3 (Static Hosting)

1. Log in to the AWS Management Console.
2. Navigate to S3 and create a new bucket.
3. Upload the HTML, and JavaScript files to the bucket.
4. Enable the static hosting option for the bucket.
5. Configure the bucket to be publicly accessible so that the application can be accessed via a URL.

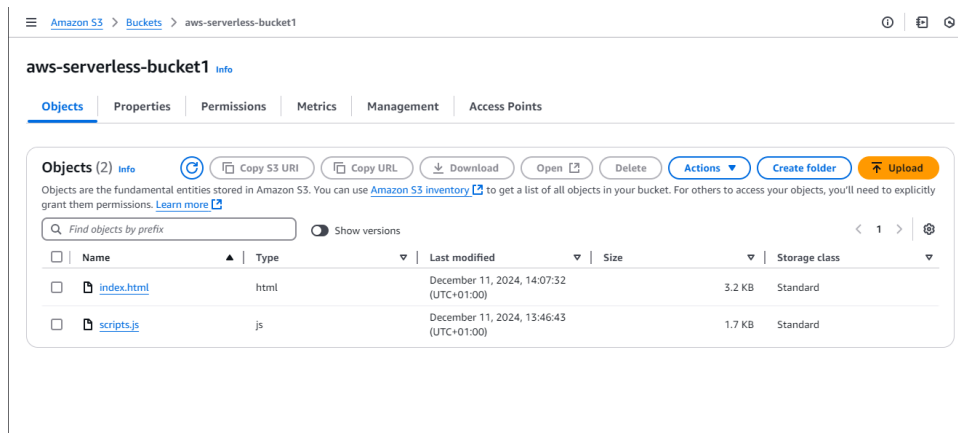


Figure 5: Bucket S3

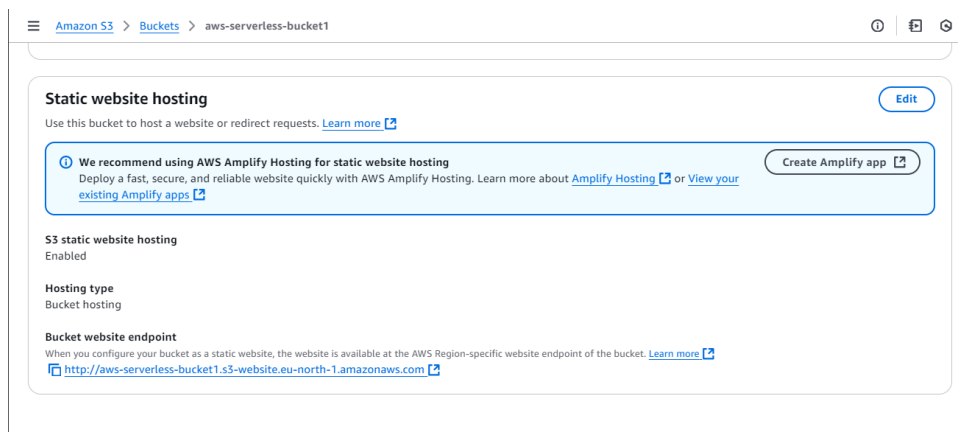


Figure 6: Enable Web Hosting in Bucket S3

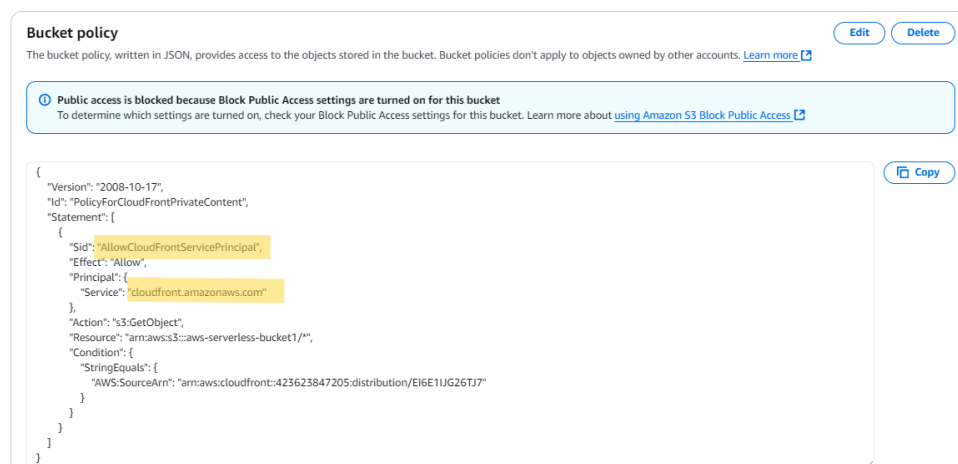


Figure 7: Bucket Policy

3.3.2 3.2 Steps to Configure CloudFront

1. Log in to the AWS Management Console and navigate to CloudFront.

2. Click on **Create Distribution**.
3. Under the **Origin Settings**, select the S3 bucket created earlier as the origin.
4. Enable **Origin Access Control (OAC)** to secure the connection between CloudFront and the S3 bucket.
5. Specify the default root object, such as `index.html`, so CloudFront knows which file to serve.
6. Modify the S3 bucket policy to allow CloudFront access by including the CloudFront service principal (`cloudfront.amazonaws.com`) in the policy.
7. Block all public access to the S3 bucket to prevent direct exposure.
8. Click on **Create Distribution** and wait for the distribution to deploy (this may take several minutes).
9. Once the deployment is complete, use the CloudFront-provided secure **HTTPS URL** to access your application.

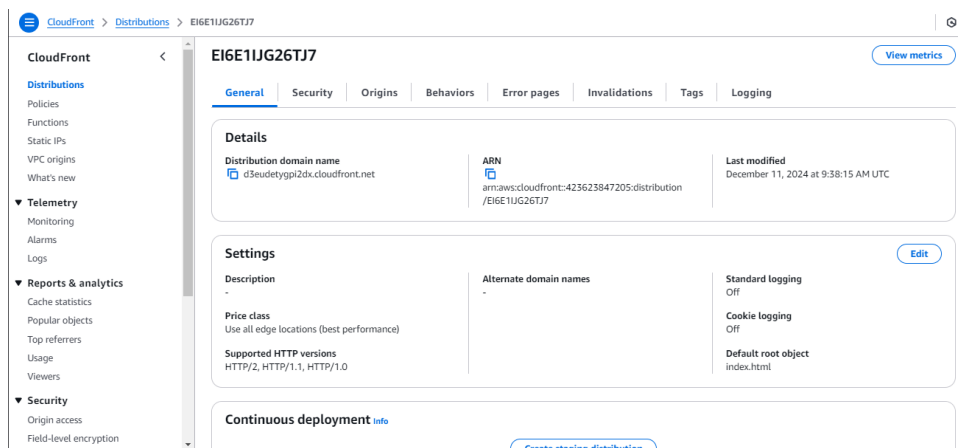


Figure 8: cloudfront distribution

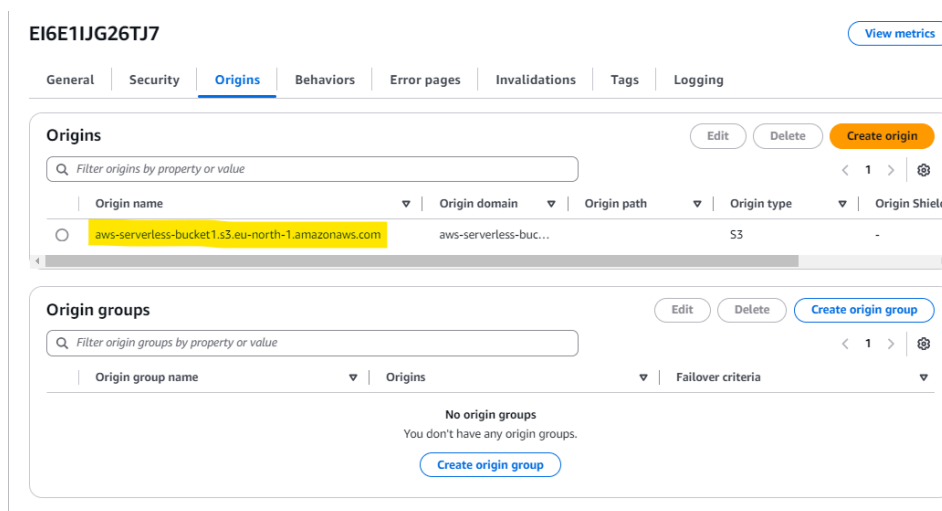


Figure 9: resource origins

3.3.3 3.3 Creating APIs with API Gateway

1. Navigate to API Gateway and create a new REST API.
2. Define resources and methods (GET, POST) to interact with the Lambda functions.
3. Link each API method to its corresponding Lambda function.
4. Configure permissions so that only authorized requests can access the API.

The screenshot shows the 'Create REST API' page in the AWS API Gateway console. The breadcrumb navigation is 'API Gateway > APIs > Create API > Create REST API'. The page title is 'Create REST API'. Under 'API details', there are three radio buttons: 'New API' (selected), 'Clone existing API', and 'Import API'. Below this, there is a text input for 'API name' with the value 'student'. There is also a 'Description - optional' text area. A dropdown menu for 'API type' is open, showing 'Regional' (selected), 'Edge-optimized', and 'Private'. A note indicates that 'Edge-optimized' and 'Private' are only accessible from VPCs.

Figure 10: Create Rest API

The screenshot shows the 'Resources' page in the AWS API Gateway console. The breadcrumb navigation is 'API Gateway > APIs > Resources - student (9r3dkbxrs2)'. The page title is 'Resources'. On the left, there is a 'Create resource' button and a list of resources: '/', 'GET', 'OPTIONS', and 'POST'. On the right, there is a 'Resource details' section for the '/' resource. It shows the 'Path' as '/' and the 'Resource ID' as 'jntd9nv15l'. There are buttons for 'API actions', 'Deploy API', 'Update documentation', and 'Enable CORS'. Below this, there is a 'Methods (3)' section with a table of methods.

| Method type | Integration type | Authorization | API key |
|-------------|------------------|---------------|--------------|
| GET | Lambda | None | Not required |
| OPTIONS | Mock | None | Not required |
| POST | Lambda | None | Not required |

Figure 11: API Resources

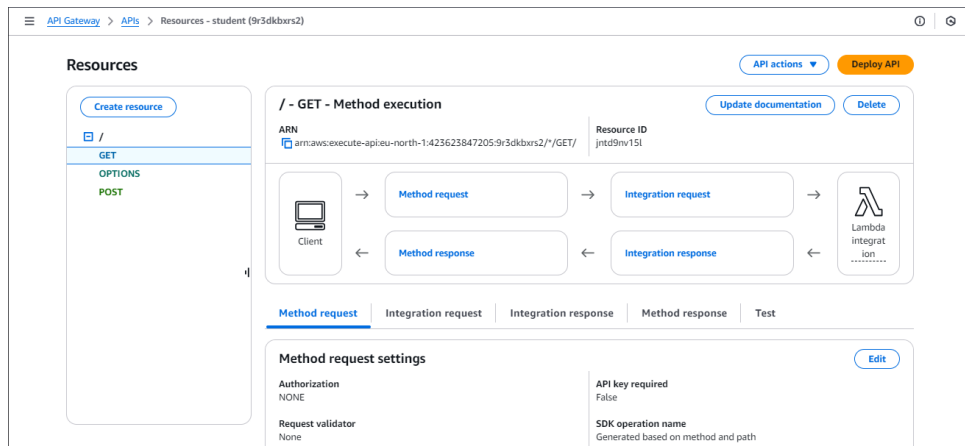


Figure 12: Example Get Resource

3.3.4 3.4 Deploying Lambda Functions

1. Go to AWS Lambda and create two functions: one for GET and one for POST.
2. Add the necessary permissions for Lambda functions to access DynamoDB.
3. Deploy the Lambda functions to make them accessible via API Gateway.

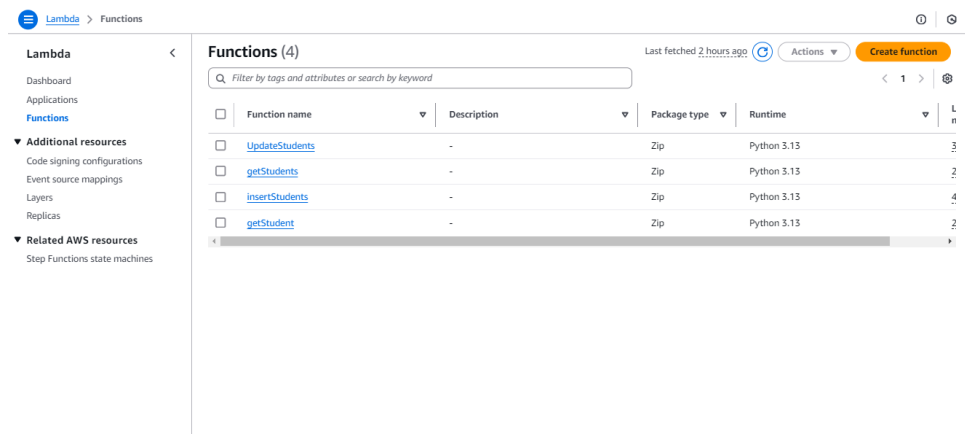


Figure 13: lambda funtions

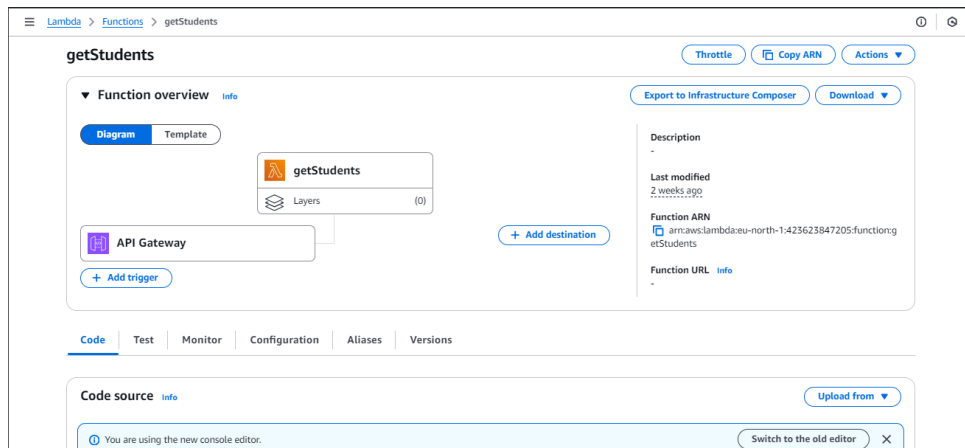


Figure 14: GET FUCNTION

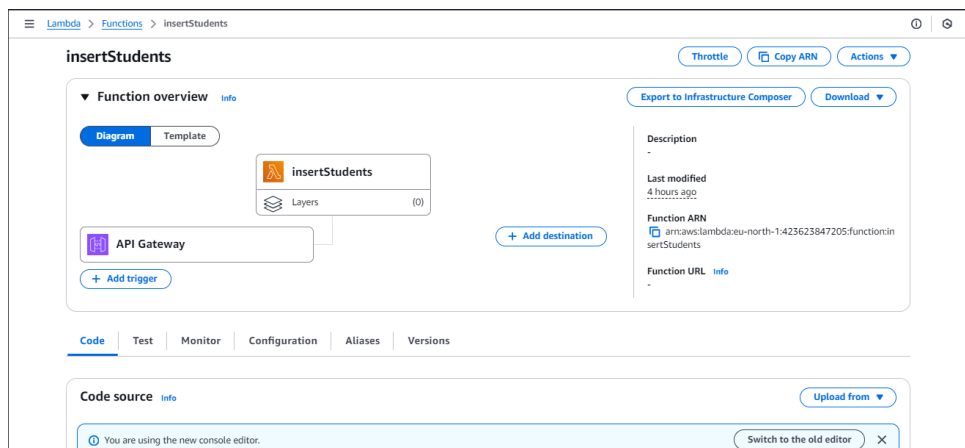


Figure 15: POST FUNCTION

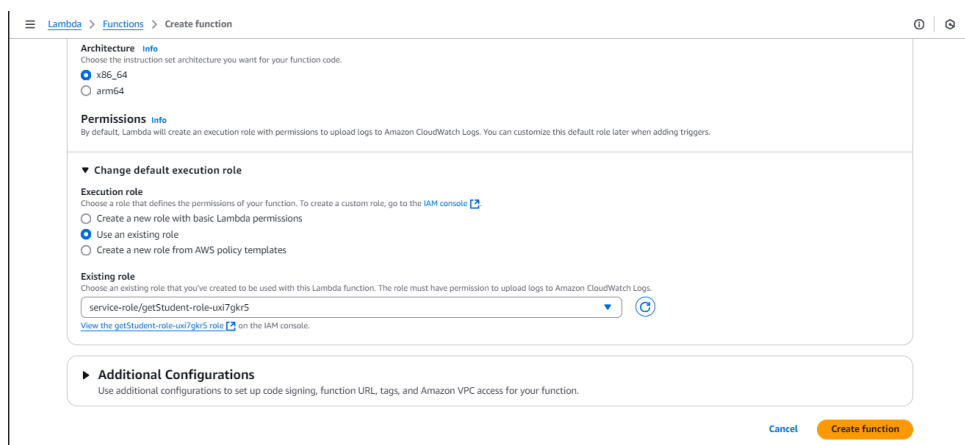


Figure 16: Permissions Lambda

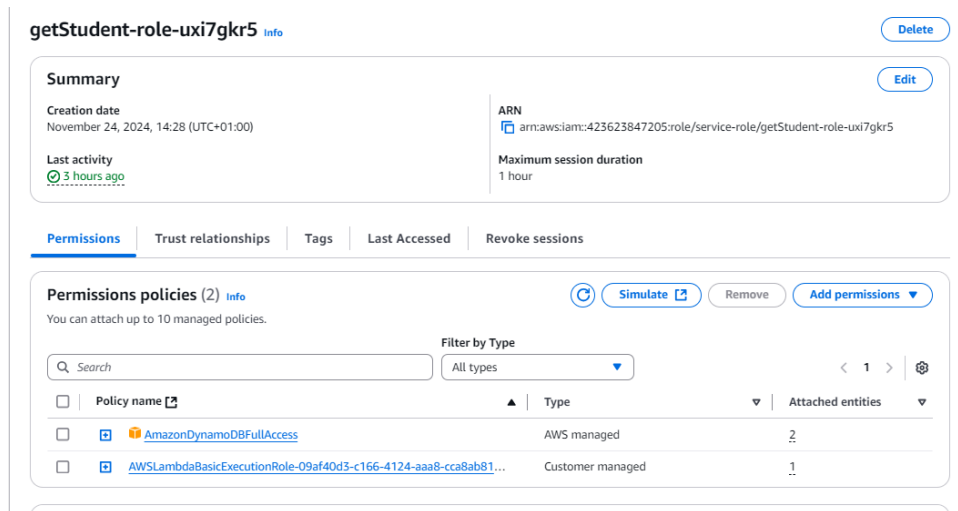


Figure 17: IAM ROLE

3.3.5 3.5 Configuring DynamoDB (Table and Data Schema)

1. Create a table in DynamoDB with `student_id` as the primary key.
2. Define a flexible schema to store student information (age , class).

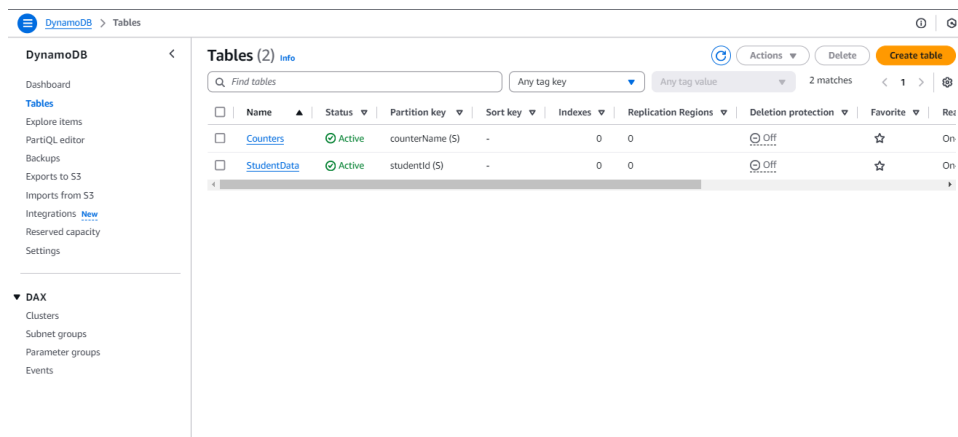


Figure 18: Dynamo DB tables

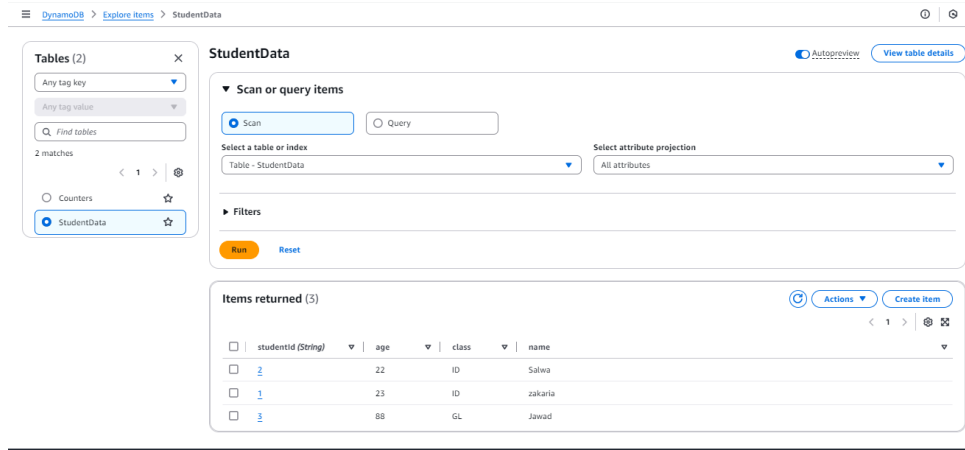


Figure 19: Explore Items

4 Data Flow

This section describes the data flow from the user submitting a request to storing it in DynamoDB.

4.0.1 4.1 Step-by-Step Explanation of a User Request

When a user submits a form to add a student:

1. **Request Submission:** The user submits a form on the frontend (delivered via CloudFront, which serves the static website from the S3 bucket).
2. **API Request Sent:** The frontend sends an HTTP POST request to the API Gateway with the student's data.
3. **Processing by Lambda:** API Gateway invokes the POST Lambda function.
4. **Storage in DynamoDB:** The Lambda function saves the data in DynamoDB.
5. **Response to User:** Lambda sends a response to the API Gateway, which relays it to the frontend, confirming successful data storage. The response is displayed securely to the user via CloudFront.

5 Project Benefits and Limitations

5.1 4.1 Benefits

The project offers several notable advantages that contribute to its efficiency, cost-effectiveness, and ease of use.

- **Low Cost with Serverless Architecture:** Using a serverless architecture with AWS Lambda, DynamoDB, and API Gateway significantly reduces infrastructure costs. There is no need to manage servers or purchase expensive hardware resources. The pay-as-you-go model ensures we only pay for the resources actually used, making the project more affordable.

- **Automatic Scalability:** The serverless architecture enables automatic scalability. As the number of users grows, AWS automatically adjusts the capacity of services like Lambda and DynamoDB to handle the load without requiring manual intervention. This scalability ensures optimal performance even during sudden traffic spikes.
- **Ease of Deployment:** Tools like AWS Lambda and API Gateway enable quick and easy deployment. Code is deployed directly to Lambda, and APIs are configured and managed via API Gateway. This approach simplifies the production process and reduces deployment-related errors.

5.2 4.2 Limitations

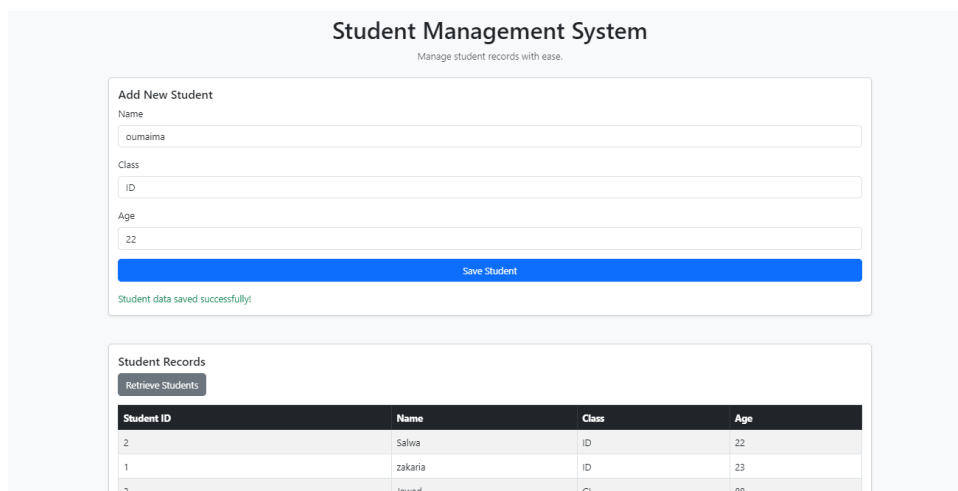
Despite its many benefits, the project has certain limitations that need consideration.

- **Dependency on AWS:** The exclusive use of AWS services like Lambda and DynamoDB creates a strong dependency on Amazon's platform. Technical issues or cost increases from AWS could directly impact the project's availability and profitability. Additionally, this dependency limits flexibility in choosing tools and providers.
- **Possible API Call Latency:** Although AWS Lambda and API Gateway offer high performance, latency can sometimes occur during API calls, especially if the Lambda service hasn't been invoked recently (cold start). This latency may affect user experience, particularly in applications requiring real-time responses.

6 Results

The project outcomes include the setup of a functional and efficient application to manage information using a serverless architecture.

- **Screenshots of the Website:** Below are screenshots of the functional website, showcasing the user interface and application usability.



Student Management System
Manage student records with ease.

Add New Student

Name
oumaima

Class
ID

Age
22

Save Student

Student data saved successfully!

Student Records

Retrieve Students

| Student ID | Name | Class | Age |
|------------|---------|-------|-----|
| 2 | Salwa | ID | 22 |
| 1 | zakaria | ID | 23 |
| 3 | Jawad | GL | 88 |

Figure 20: Screenshot of the app

- **Example of Data Stored in DynamoDB:** Here is an example of the data stored in the DynamoDB table for a student.

```
{
  "student_id": "12345",
  "name": "John Doe",
  "class": "Math",
  "Age": "23"
}
```

- **Tests Conducted to Verify Functionality:** Tests were conducted to validate the functionality of the website and APIs. These tests included data submissions via the form and verification of data storage in DynamoDB. All tests showed that data was correctly recorded and retrieved.

7 Conclusion and Future Work

This project demonstrated the effectiveness of serverless architecture for managing student information. By leveraging AWS services such as Lambda, DynamoDB, and API Gateway, the application proved to be performant, scalable, and cost-efficient.

- **Project Summary and Results:** The project successfully implemented a system that enables submitting and retrieving information via a static frontend website hosted on S3 and backend AWS Lambda functions. Data was correctly processed and stored in DynamoDB, and the application showcased the advantages of a serverless environment.
- **Suggestions for Improving the Project:**
 - **Add Authentication:** Integrate AWS Cognito to ensure only authorized users can submit or retrieve information, improving security.
 - **Enhance Logging and Monitoring:** Use AWS CloudWatch to monitor application performance and troubleshoot issues efficiently.
 - **Implement Version Control:** Use AWS CodePipeline to automate deployment and versioning of the application for continuous delivery.
 - **Data Backup and Recovery:** Use AWS Backup to automate periodic backups of DynamoDB and ensure data durability.
 - **Integration with AI Services:** Incorporate AWS SageMaker for predictive analytics or personalization features, such as suggesting courses based on student data.
 - **Event-Driven Architecture:** Use AWS EventBridge to trigger actions based on specific events, such as sending notifications for new data entries.
 - **Compliance and Security Enhancements:** Implement AWS Identity and Access Management (IAM) policies and AWS Shield to secure the application against DDoS attacks.
 - **Cost Optimization:** Use AWS Cost Explorer and Trusted Advisor to identify areas to further optimize resource usage and minimize costs.

- **Serverless CI/CD Pipeline:** Implement a fully serverless CI/CD pipeline using AWS CodeBuild and AWS CodeDeploy.